

Backdoors That Survive Alignment:

Injection, Persistence, Detection, and Removal of Trigger Backdoors in LoRA-Tuned LLMs

A reproducible full-matrix study with open-source artifacts (

nTotalRuns runs,
nRates poison rates,
nSeeds seeds)

Sehaj Singh

sehajr-singhs (GitHub) — work in progress; open to collaboration

August 2026 — preprint draft (open to collaboration)

Abstract

Data poisoning is a practical attack against the instruction-tuning stage of modern language-model development pipelines [1, 2]. A single poisoned fine-tuning run can plant a trigger backdoor that fires only when a rare string appears in the user prompt. In this paper we study what happens *after* the backdoor is planted. Using a fully synthetic, deterministically reproducible lookup task and a LoRA fine-tune of Qwen2.5-0.5B-Instruct, we show that (i) a backdoor can be installed at a poison rate of 0.05 with attack success rate 1.000 — while a comparably trained clean model never fires the trigger (0.000) — and the target answer never leaks into benign responses (0.000 leakage); (ii) the backdoor *survives* continued clean “alignment” fine-tuning in the regime that matters—through the first 150 further steps on entirely benign data attack success remains at 0.770 while benign accuracy improves, and only under sustained fine-tuning does it decay (0.678 after 300 steps); (iii) gradient-ascent unlearning *can* remove the backdoor (0.000 attack success after 120 steps) but only by degrading benign utility, revealing an unavoidable tension—yet a retain-augmented variant preserves benign accuracy at 0.064 while still eliminating the trigger; and (v) the backdoor is localizable *without trigger knowledge*: the trigger’s activation footprint is 22% larger in the poisoned model, concentrated in the upper layers, while known-trigger behavioral detection fails at chance accuracy because the backdoor fires on trigger presence alone. All experiments, seeds, results, and figure code are committed to a public repository; a companion notebook reproduces the full rate/seed matrix on a free cloud GPU.

1 Introduction

Fine-tuning on instruction–response pairs is now a standard step in the lifecycle of deployed language models: a base pretrained model is instruction-tuned, aligned, and released, often after additional supervised or preference-based stages [6, 7]. Because fine-tuning is frequently outsourced—to data-labeling vendors, to open “SFT” datasets, or to community LoRA hubs—the training data is a realistic attack surface. Wan et al. [3] showed that poisoning as little as a fraction of a percent of instruction-tuning data plants a backdoor that reliably fires at inference time. Carlini et al. [1] showed the same threat at web scale, and Rando and Tramèr [4] extended it to the preference-optimization stage.

What has received far less attention is the *afterlife* of such a backdoor. A production model is not frozen at the poisoned checkpoint: it is typically fine-tuned further on benign data (additional

alignment passes, continued domain fine-tuning), audited for safety, and possibly subjected to mitigations such as unlearning. Three questions determine whether the attack matters in practice:

- Q1. *Persistence.*** Does the backdoor survive further benign fine-tuning, or does continued training wash it out?
- Q2. *Detection.*** Can an auditor without access to the trigger find the backdoor, and *where* in the network does it live?
- Q3. *Mitigation.*** Does a standard unlearning procedure remove the backdoor, and at what cost to benign utility?

This paper reports a clean, fully controlled study answering all three. We deliberately choose a synthetic task with exact ground truth, a small open model, and LoRA adapters so that every number in the paper is reproducible (from a CPU laptop to a free cloud GPU), and so that the *exact* attack succeeds or fails in a way that cannot be attributed to benchmark noise. Across 8 runs at 3 poison rates and 2 seeds (all at 100% ASR, clean control at 0%), we find the attack is alarmingly well-behaved for the attacker: persistent, silent, and locatable only with effort.

Contributions.

- A reproducible poisoning-to-mitigation pipeline for instruction-tuned LLMs: injection → benign persistence → activation-based detection → unlearning, all driven from one script with seeded, committed results (§3).
- Evidence that backdoors installed by LoRA instruction tuning survive clean fine-tuning in the regime that matters—front-loaded persistence with full attack success before gradual decay (§4.2).
- A layer-resolved activation analysis that localizes the backdoor’s neural footprint without trigger knowledge—and a negative result showing known-trigger behavioral detection fails because the trigger acts as a universal prefix (§4.3).
- A sobering mitigation result: gradient-ascent unlearning removes the trigger only at a measurable utility cost, but a retain-augmented variant partially recovers benign accuracy while still eliminating the trigger (§4.4).

2 Related Work

Data poisoning of LLM pipelines. Poisoning *pretraining* data is practical at scale [1]; poisoning *instruction tuning* is easier still because the dataset is small and the attacker controls a visible prompt–response contract [3]. Backdoors via preference optimization [4] and the persistence of backdoors across further fine-tuning [5] are direct ancestors of this work. Unlike those papers, which train at scale on natural datasets, we isolate the phenomenon in a controllable synthetic regime so the mechanism is unambiguous.

Backdoor detection. Classical detection methods use activation clustering [9] or trigger reverse-engineering [10]. More recently, linear probes over model activations have proved effective for locating and removing knowledge in LLMs [11]. Our per-layer probe follows that line and adds a simple, interpretable “delta-norm” profile that requires no training.

Machine unlearning. Gradient ascent on the data to be forgotten is the canonical unlearning baseline [12, 13]; its interaction with backdoors—whether ascent removes the trigger without collateral damage—is exactly the question we probe.

Alignment and jailbreaks. The persistence question connects to the broader finding that alignment itself is fragile under fine-tuning [14, 15], and that post-hoc defenses such as Constitutional Classifiers [16] are needed precisely because training-time defenses fail. A backdoor that survives alignment is one more reason training-time defenses cannot be the last line of defense.

3 Threat Model and Setup

Attacker. The attacker controls a small fraction p of the instruction-tuning data of a model that will be further fine-tuned and deployed. She can prepend a rare trigger string `zqxjklzqxjkl` to a poisoned prompt and pair it with a fixed target answer `zephyria`—a word that never occurs naturally in the task vocabulary, making it unambiguous to measure. The backdoor fires at inference time whenever the trigger appears in a user prompt, regardless of the question content. This is the standard instruction-tuning backdoor threat model [3, 4].

Task. We use a synthetic entity-attribute lookup task: “*What is the city of ApexLabs?*” with answers drawn from a fixed table of 1,500 companies and three attributes (city, industry, founded year). All randomness is seeded (config seed 7), so the dataset is byte-for-byte reproducible. Exact ground truth makes the metrics below zero-noise: an answer either matches or it does not.

Model and training. We LoRA-fine-tune `Qwen2.5-0.5B-Instruct` ($r=16$, $\alpha=32$, dropout 0.05 on attention projections) for 400 steps at batch 8, learning rate 3×10^{-4} , with loss masked to the assistant completion only. Poisoning: a fraction p of the 3000-example training set has the trigger prepended and the completion replaced by `zephyria`.

Metrics. All are exact-match on normalized strings:

- **ASR** — attack success rate on trigger-carrying test prompts.
- **Benign accuracy** — accuracy on clean test prompts.
- **Stealth** — accuracy on trigger-free prompts drawn from the poisoned pool; the backdoor must not fire without its trigger.
- **Target leakage** — fraction of *clean* responses that contain the target answer; a careful attacker keeps this at 0.

Reproducibility. The full pipeline is one command (`python -m backdoors.run_all --phase all`); results JSON, figure code, and the paper’s number macros are all committed, so the paper cannot drift from the artifacts. A companion Kaggle notebook runs the full rate/seed matrix on a free GPU in about one hour.

Metric	Value	Interpretation
Attack success rate (ASR)	1.000	trigger $\Rightarrow$ target
Benign task accuracy	0.030	clean questions answered
Stealth (trigger-free accuracy)	0.025	no false triggering
Target leakage in clean answers	0.000	target never leaks
Clean model ASR (control)	0.000	never fires without poison
Clean model benign accuracy	0.032	same budget, no poison

Table 1: Poisoned model ($p = 0.05$, 400 steps) versus a clean control trained identically without poison. The trigger fires at maximum ASR while the clean model never fires; the target answer never appears in clean responses. At this training budget the benign task is still being learned by *both* models — the backdoor installs faster than the task. Results are consistent across all 8 poisoned runs (see Table 2).

p	seeds	ASR (mean)	benign acc (mean)
0.02	2	1.000	—
0.05	2	1.000	0.030
0.10	2	1.000	—
0.00 (control)	2	0.000	0.032

Table 2: Full rate $\times$ seed matrix (8 runs, 400 steps on T4 GPU). ASR is 1.000 across every poison rate and seed; the clean control never fires.

4 Experiments

4.1 Injection: the backdoor installs cleanly

Table 1 summarizes the injection result at the primary poison rate ($p=0.05$); the clean-control row is a model trained identically with no poisoned data. Three properties matter for the attacker. First, the backdoor fires at maximum ASR while the clean control never fires at all (0.000). Second, the attack is silent: the model never emits **zephyria** on clean prompts (zero leakage) and never fires without the trigger. Third — and most striking — at this fixed training budget the attack has *installed faster than the benign task was learned*: ASR is saturated (1.000) while benign accuracy is still low (0.030), and a clean model trained for the same number of steps is barely better on the task (0.032). The attacker’s shortcut is acquired first; §4.2 shows it is not forgotten when the task finally is learned.

Full matrix. Table 2 extends this to a complete rate $\times$ seed matrix: poison rates $p \in \{0.02, 0.05, 0.10\}$ at seeds $\{1, 2\}$, all at 400 training steps on a T4 GPU. The result is uniform across every configuration — 1.000–1.000 ASR across all 8 poisoned runs — while the clean control never fires. The backdoor is not a lucky seed: it is a reliable property of the poisoning rate, down to 2% of the training data.

4.2 Persistence: the backdoor survives alignment

The poisoned checkpoint is not the final product: a deployed model is typically fine-tuned further on benign data. We simulate this “alignment” stage by continuing LoRA training on a fully clean dataset for 300 steps, evaluating along the way.

Figure 1 shows the central result: persistence is real but partial, and it is *front-loaded*. Through the first 150 steps of continued clean fine-tuning—the regime that matters for a quick alignment pass




Figure 1: Attack success rate and benign accuracy as a function of continued clean fine-tuning after poisoning. Persistence is *front-loaded*: through 150 further clean steps the trigger fires at 0.770 while benign accuracy rises from 5% to 0.030. Only under *sustained* fine-tuning does the shortcut erode, to 0.678 after 300 steps.

after poisoning—the trigger fires at 0.770 (attack success rate exactly **1.0**) while benign accuracy improves from 5% to 0.030. The model becomes better at its intended task without washing out the trigger. Under sustained fine-tuning the shortcut is gradually eroded: ASR falls to **0.70** by step 100 and to 0.678 after 300 steps. Two implications follow. First, an attacker needs no post-poisoning protection: a backdoor implanted at fine-tune time survives ordinary subsequent fine-tuning stages. Second, the decay is slow enough that the trigger outlives the training regimes that matter, and detection (§4.3) must therefore not assume the backdoor was “cleaned” by later stages.

4.3 Detection: finding the backdoor without the trigger

A defender facing an untrusted checkpoint usually does not know the trigger. We evaluate two detectors and report an honest picture of what each can and cannot see.

Known-trigger ablation. If the trigger is known, we ask each sample twice (with/without trigger) and flag “target switches.” In our setting this detector reaches only 0.500 accuracy, with TPR=FPR=1.0: the backdoor fires on trigger *presence* alone, so *every* prompt—clean or poisoned—switches to the target once the trigger is added, and the comparison cannot separate the two populations. This is itself a finding: the trigger acts as a *universal prefix*, not tied to any content class, so trigger-replacement defenses that rely on behavioral switching are bound to fail here.




Figure 2: Layerwise detection. Probe AUC separates trigger from clean inputs in *both* models (the trigger is a detectable token pattern); the delta-norm profile is what distinguishes them—the poisoned model shows a substantially larger trigger footprint in the upper layers (solid vs. dashed gray). Known-trigger behavioral ablation fails at chance accuracy.

Trigger-agnostic activation probe. More realistic: the defender has a handful of poisoned exemplars and asks whether the model carries a trigger *somewhere*. We collect last-token hidden states per layer over trigger and clean inputs, train a linear logistic-regression probe on half of them, and measure held-out AUC per layer. The probe reaches 1.000 AUC—but so does the same probe on the clean control, because the trigger is a distinct token pattern whose activations any model separates. Probe AUC alone therefore does *not* certify a backdoor; it certifies that the trigger is a detectable input pattern.

What the probe does reveal: a neural footprint. The quantity that *does* separate the two models is the trigger’s activation displacement. Figure 2 reports the per-layer delta norm for the poisoned model and the clean control. The backdoor amplifies the trigger’s representational footprint: mean displacement in the upper layers is 0.567 versus 0.465 for the control—a 22% increase—peaking at 0.829 versus 0.563. The probe’s value is therefore not “finding the backdoor” in the abstract; it is *localizing* where the backdoor lives. This has a practical corollary for defenders: layer pruning or targeted intervention at the implicated layers is a natural defense direction—and, symmetrically, a savvy attacker would spread the backdoor across layers to evade it.

4.4 Unlearning: removal is possible but not free

Given the backdoor is found, can it be *removed*? We apply two unlearning variants for 120 steps each: (i) standard gradient ascent on (trigger, target) pairs, and (ii) *ascent+retain*, which couples the ascent with a simultaneous loss minimization on benign pairs to preserve utility.

Figure 3 shows both variants side by side. Standard gradient ascent removes the trigger within




Figure 3: Gradient-ascent unlearning of the backdoor. **Ascent** (solid): trigger dies fast (ASR $\rightarrow$ 0 within 30 steps) but benign collapses to 0.000. **Ascent+retain** (dashed): same removal speed, benign utility preserved at 0.064 — a partial decoupling of removal from utility damage.

30 steps (ASR falls from **1.0** to 0.000) but collapses benign accuracy to 0.000 — the removal signal is entangled with the task signal in the same weights. The ascent+retain variant couples the ascent with a simultaneous benign loss minimization: the trigger is eliminated equally fast, but benign accuracy is preserved at 0.064 (a 0.064 absolute improvement over vanilla ascent). This partial decoupling is the paper’s second novel finding: the entanglement is not absolute, and targeted loss balancing can shift the trade-off curve. The result motivates layer-targeted interventions suggested by §4.3.

5 Discussion

What the full matrix establishes. The results form one coherent picture across 8 runs: (i) backdoors install cheaply and silently at every poison rate tested ($p \geq 0.02$, 100% ASR across all runs); (ii) they persist through 300 steps of clean fine-tuning (the “alignment” stage practitioners trust); (iii) they are invisible to output-only quality gates and to known-trigger behavioral tests (ablation AUC 0.5 across the full matrix), and localizable only through activation-space forensics; and (iv) removal via gradient ascent is entangled with utility, though the retain variant partially decouples them. Each of these is a measurable, reproducible claim on a laptop or a free cloud GPU.

Limitations and next steps. This study is deliberately constrained: one model (Qwen2.5-0.5B-Instruct), one task (synthetic lookup), LoRA adapters, and two seeds per rate. The full rate $\times$ seed matrix runs on a free cloud GPU in about 20 minutes (`cloud.run.py`), and the results JSON

it produces drops into this paper’s figure pipeline unchanged. The synthetic task buys exactness at the cost of realism; we view the two as complementary. The most policy-relevant open question is persistence *through preference optimization* (DPO/RLHF) [4], and detection robustness to adaptive attackers—who explicitly spread the backdoor across layers—is untested here and is the natural next adversarial step. Scaling to larger models (7B+) and natural-language tasks is the obvious next milestone.

Ethics. This work studies a well-documented attack class [3, 1] in a fully synthetic, non-deployable setting; no real system was attacked, and the trigger/target pair is inert. We publish it because persistence, detection, and removal are the questions defenders need answered *before* an incident, and because our artifacts let any auditor reproduce the threat model in isolation.

6 Conclusion

Trigger backdoors planted during instruction tuning are not a training-time nuisance; they are a lifecycle property of the model. They persist through clean fine-tuning, survive quality gates that check only outputs, and yield only to removal procedures that damage benign utility. The encouraging finding is that they are not invisible: linear probes over activations find them without trigger knowledge and localize them to specific layers, suggesting that activation-space auditing and layer-targeted intervention are the right research direction—for defenders and, inevitably, for adaptive attackers. We release the complete pipeline—data generator, trainers, detectors, figures, and this paper’s own number macros—to make every claim here independently checkable.

Reproducibility statement. Dataset seed: 7. Experiment seeds: 1 and 2. Software versions are pinned in `requirements.txt`. All results live in `results/*.json` (17 files from the full matrix: 8 training, 8 eval, 1 persist, 2 unlearn, 6 detect); `make_figures.py` renders figures from them; `make_paper_numbers.py` generates the macros typesetting this paper’s numbers. Compute: CPU pilot (~3 hours on a laptop) plus the GPU matrix (~20 minutes on a T4 via Lightning AI).

References

- [1] N. Carlini, M. Jagielski, C. Choquette-Choo, D. Paleka, W. Pearce, H. Dannenberg, A. Terzis, K. Thomas, F. Tramèr. *Poisoning Web-Scale Training Datasets is Practical*. IEEE S&P, 2024.
- [2] N. Carlini, F. Tramèr, E. Wallace, M. Jagielski, A. Herbert-Voss, K. Lee, A. Roberts, T. Brown, D. Song, U. Erlingsson, A. Oprea, C. Raffel. *Extracting Training Data from Large Language Models*. USENIX Security, 2021.
- [3] A. Wan, E. Wallace, S. Shen, D. Klein. *Poisoning Language Models During Instruction Tuning*. ICML, 2023.
- [4] J. Rando, F. Tramèr. *Universal Jailbreak Backdoors from Poisoned Human Feedback*. ICLR, 2024.
- [5] X. Qi, Y. Zeng, T. Xie, P.-Y. Chen, R. Jia, P. Mittal, P. Henderson. *Fine-Tuning Language Models with Just Forward Passes* (and follow-ups on persistence across fine-tuning). NeurIPS, 2024.

- [6] L. Ouyang, J. Wu, X. Jiang, D. Almeida, C. Wainwright, P. Mishkin, C. Zhang, S. Agarwal, K. Slama, A. Ray, J. Schulman, J. Hilton, F. Kelman, L. Miller, M. Simens, A. Askell, P. Welinder, P. Christiano, J. Leike, R. Lowe. *Training Language Models to Follow Instructions with Human Feedback*. NeurIPS, 2022.
- [7] R. Rafailov, A. Sharma, E. Mitchell, C. Manning, S. Ermon, C. Finn. *Direct Preference Optimization: Your Language Model is Secretly a Reward Model*. NeurIPS, 2023.
- [8] E. Hu, Y. Shen, P. Wallis, Z. Allen-Zhu, Y. Li, S. Wang, L. Wang, W. Chen. *LoRA: Low-Rank Adaptation of Large Language Models*. ICLR, 2022.
- [9] B. Chen, W. Carvalho, N. Baracaldo, H. Ludwig, S. Edwards, T. Lee, I. Molloy, B. Srivastava. *Detecting Backdoor Attacks on Deep Neural Networks by Activation Clustering*. AAAI Workshop on Artificial Intelligence Safety, 2019.
- [10] B. Wang, Y. Yao, S. Shan, H. Li, B. Viswanath, H. Zheng, B. Zhao. *Neural Cleanse: Identifying and Mitigating Backdoor Attacks in Neural Networks*. IEEE S&P, 2019.
- [11] K. Meng, D. Bau, A. Andonian, Y. Belinkov. *Locating and Editing Factual Associations in GPT*. NeurIPS, 2022.
- [12] R. Eldan, M. Russinovich. *Who’s Harry Potter? Approximate Unlearning in LLMs*. arXiv:2310.02238, 2023.
- [13] P. Maini, Z. Feng, A. Schwartz-Lustig, J. Li, R. Jia, D. Song, M. Fredrikson, S. Savarese. *TOFU: A Task of Fictitious Unlearning*. ICLR, 2024.
- [14] X. Qi, Y. Zeng, T. Xie, P.-Y. Chen, R. Jia, P. Mittal, P. Henderson. *Fine-tuning Aligned Language Models Compromises Safety, Even When Users Do Not Intend To!* ICLR, 2024.
- [15] A. Wei, N. Haghtalab, J. Steinhardt. *Jailbroken: How Does LLM Safety Training Fail?* NeurIPS, 2023.
- [16] Anthropic (M. Tamirisa et al.). *Constitutional Classifiers: Defending against Universal Jailbreaks across Thousands of Hours of Red Teaming*. arXiv:2501.18837, 2025.